



Assignment Solution

Week12: Apache Spark - Structured API
Part-2

Spark Structured APIs - Assignment Solutions

Assignment 1 :

Given 2 Datasets employee.json and dept.json

We need to calculate the count of employees against each department. Use Structured APIs.

Code:

//Find the count of employees against each department

```
import org.apache.spark.SparkConf
import org.apache.spark.sql.SparkSession
import org.apache.log4j.Level
import org.apache.log4j.Logger
import org.apache.spark.sql.functions._

object Assignment1_Week12 extends App{

  //Setting the Log Level
  Logger.getLogger("org").setLevel(Level.ERROR)

  //Setting the spark conf
  val sparkConf = new SparkConf()
  sparkConf.set("spark.app.name","Assignment1_Week12")
  sparkConf.set("spark.master","local[2]")

  //Creating Spark Session
  val spark = SparkSession.builder()
    .config(sparkConf)
    .getOrCreate()

  //Load the department data into a Dataframe using dataframe reader API

  val deptDf = spark.read
    .format("json")
    .option("path","C:/TrendyTech/SparkExamples/dept.json")
    .load()

  // deptDf.show()
  // deptDf.printSchema()
```

//Load the employee data into a Dataframe using dataframe reader API

```
val employeeDf = spark.read
  .format("json")
  .option("path","C:/TrendyTech/SparkExamples/employee.json")
  .load()
```

```
// employeeDf.show()
// employeeDf.printSchema()
```

//Joining of two dataframes using left outer join, with department dataframe on left side

```
val joinCondition = deptDf.col("deptid") === employeeDf.col("deptid")//join
condition
```

```
val joinType = "left" //joinType
```

```
val joinedDf = deptDf.join(employeeDf, joinCondition, joinType) //Joining of two
dataframes
```

//drop the ambiguous column deptid of employee dataframe,from the joined Dataframe

```
val joinedDfNew = joinedDf.drop(employeeDf.col("deptid"))
```

//Use first function so as to get other columns also along with aggregated columns

```
joinedDfNew.groupBy("deptid").agg(count("empname").as("empcount"),first("deptName").as("deptName")).dropDuplicates("deptName").show()
```

```
spark.stop()
}
```

Output:

deptid	empcount	deptName
21	1	HR
41	2	Fin
51	0	Admin
31	1	Marketing
11	1	IT

Assignment 2

Find the top movies as shown in spark practical 18 using broadcast join. Use Dataframes or Datasets to solve it this time.

Code:

```
import org.apache.spark.SparkConf
import org.apache.spark.sql.SparkSession
import org.apache.log4j.Level
import org.apache.log4j.Logger
import org.apache.spark.sql.functions._

object Assignment2_Week12 extends App {

  //Setting the Log Level
  Logger.getLogger("org").setLevel(Level.ERROR)

  //Setting the spark conf
  val sparkConf = new SparkConf()
  sparkConf.set("spark.app.name","Assignment2_Week12")
  sparkConf.set("spark.master","local[2]")

  //Creating Spark Session
  val spark = SparkSession.builder()
    .config(sparkConf)
    .getOrCreate()

  //Creation of a ratings dataframe using a case class approach

  case class Ratings(userid:Int,movieid:Int,rating:Int,timestamp:String)//create a
  case-class that represents the schema

  //Creation of base RDD for ratings data
  val ratingsRDD =
  spark.sparkContext.textFile("C:/TrendyTech/SparkExamples/ratings.dat");//ratings data
  does not have a schema, so first loading to an RDD

  // map the RDD elements into instances of the case class

  val caseClassSchemaRDD = ratingsRDD.map(x => x.split(":")).map(x =>
  Ratings(x(0).toInt,x(1).toInt,x(2).toInt,x(3)) )
```

//Transform to a Dataframe:

```
import spark.implicits._
```

```
val ratingsDf = caseClassSchemaRDD.toDF()
```

```
// ratingsDf.show()
```

```
// ratingsDf.printSchema()
```

//Creation of base RDD for movies data

```
val moviesRDD =
```

```
spark.sparkContext.textFile("C:/TrendyTech/SparkExamples/movies.dat")
```

//defining the schema using case class

```
case class Movies(movieid:Int,moviename:String,genre:String)
```

```
val moviestransformedRDD = moviesRDD.map(line => line.split(":")).map(fields =>
Movies(fields(0).toInt,fields(1),fields(2)) )
```

```
val moviesNewDf =
```

```
moviestransformedRDD.toDF().select("movieid","moviename")
```

```
// moviesNewDf.show()
```

```
//moviesNewDf.printSchema()
```

```
val transformedmovieDf = ratingsDf.groupBy("movieid")
```

```
.agg(count("rating").as("movieViewCount"),avg("rating").as("avgMovieRating"))
```

```
.orderBy(desc("movieViewCount"))
```

```
//transformedmovieDf.show()
```

```
val popularMoviesDf = transformedmovieDf.filter("movieViewCount > 1000 AND
avgMovieRating > 4.5")
```

```
// popularMoviesDf.show()
```

//Now we want to associate the Movie names also, so we use a broadcast join

```
spark.sql("SET spark.sql.autoBroadcastJoinThreshold = -1")
```

```
val joinCondition = popularMoviesDf.col("movieid") ===  
moviesNewDf.col("movieid") //join condition
```

```
val joinType = "inner" //type of  
join
```

```
val finalPopularMoviesDf =  
popularMoviesDf.join(broadcast(moviesNewDf),joinCondition,joinType).drop(popularM  
oviesDf.col("movieid")).sort(desc("avgMovieRating")) //joining the 2 dataframes using  
broadcast join where movies data is the smaller dataset
```

```
finalPopularMoviesDf.drop("movieViewCount","movieid","avgMovieRating").show(false  
)
```

```
spark.stop()
```

```
}
```



Output:

moviename
Shawshank Redemption, The (1994)
Godfather, The (1972)
Usual Suspects, The (1995)
Schindler's List (1993)

Assignment 3

File A is a text file of size 1.2 GB in HDFS at location /loc/x. It contains match by match statistics of runs scored by all the batsman in the history of cricket.

File B is a text file of size 1.2 MB present in local dir /loc/y. It contains list of batsman playing in cricket world cup 2019.

File A:

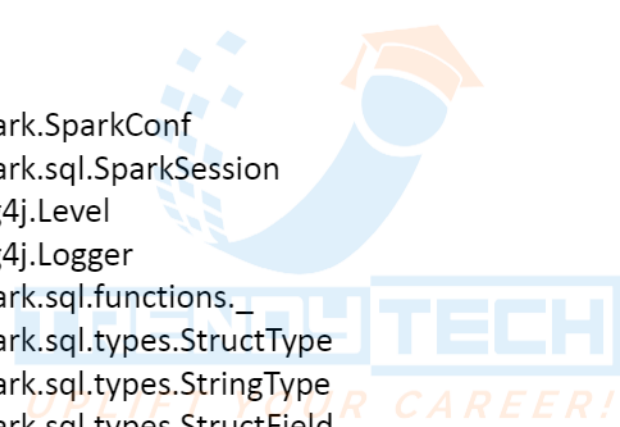
1 Rohit_Sharma India 200 100.2
 1 Virat_Kohli India 100 98.02
 1 Steven_Smith Aus 77 79.23
 35 Clive_Lloyd WI 29 37.00
 243 Rohit_Sharma India 23 150.00
 243 Faf_du_Plessis SA 17 35.06

File B:

Rohit_Sharma India
 Steven_Smith Aus
 Virat_Kohli India

Find the batsman participating in 2019 who has the best average of scoring runs in his career. Solve this using Dataframes or Datasets.

**** File is tab separated.Headers not part of file**

Code:


```
import org.apache.spark.SparkConf
import org.apache.spark.sql.SparkSession
import org.apache.log4j.Level
import org.apache.log4j.Logger
import org.apache.spark.sql.functions._
import org.apache.spark.sql.types.StructType
import org.apache.spark.sql.types.StringType
import org.apache.spark.sql.types.StructField
import org.apache.spark.sql.Row

object Assignment3_Week12 extends App {

  //Setting the Log Level
  Logger.getLogger("org").setLevel(Level.ERROR)

  //Setting the spark conf
  val sparkConf = new SparkConf()
  sparkConf.set("spark.app.name","Assignment3_Week12")
  sparkConf.set("spark.master","local[2]")

  //Creating Spark Session
  val spark = SparkSession.builder()
    .config(sparkConf)
    .getOrCreate()
```



```
//Case class creation
```

```
case class BatsmenHistory(MatchNumber:Int,Batsman:String,Team:String,
RunsScored:Int,StrikeRate:Double)
```

```
//Creation of base RDD for historical data
```

```
val batsmenHistoryRDD =
spark.sparkContext.textFile("C:/TrendyTech/SparkExamples/FileA_BatsmenDetails_History.txt")
```

```
val batsmenHistorySchemaRDD = batsmenHistoryRDD.map(line =>
line.split("\t")).map(fields =>
BatsmenHistory(fields(0).toInt,fields(1),fields(2),fields(3).toInt,fields(4).toDouble) )
```

```
// Dataframe creation
```

```
import spark.implicits._
```

```
val batsmenHistoryDf = batsmenHistorySchemaRDD.toDF()
```

```
//batsmenHistoryDf.show()
```

```
//batsmenHistoryDf.printSchema()
```

```
//Calculating Average runs scored by a batsman in history, with highest average at top
```

```
val batsmenBestRunsAvgHistoryDf =
batsmenHistoryDf.groupBy("Batsman").agg(avg("RunsScored").as("AverageRunsScored"))
.select("Batsman","AverageRunsScored")
```

```
//batsmenBestRunsAvgHistoryDf.sort(col("AverageRunsScored").desc).show()
```

```
//create a base RDD from input data of worldcup
```

```
val batsmenWorldCupRDD =
spark.sparkContext.textFile("C:/TrendyTech/SparkExamples/FileB_BatsemenDetails_Worldcup2019.txt")
```

```
//Alternative Approach instead of using case class ,though case class can also be used
instead-
```


//Programmatically create an explicit schema of the worldcup 2019 file:

```
val batsmenworldcupSchema = StructType(List(
  StructField("batsman",StringType,false),
  StructField("team",StringType)
))
```

//Convert RDD[Array[String]] to RDD[Row].

```
val batsmenWorldCupRowRDD = batsmenWorldCupRDD.map(line =>
line.split("\t")).map( fields => Row(fields(0),fields(1)))
```

//Apply the explicitly defined Struct Type schema to the RDD[Row]

```
val batsmenWorldCupDf = spark.createDataFrame(batsmenWorldCupRowRDD,
batsmenworldcupSchema)
```

```
batsmenWorldCupDf.show()
batsmenWorldCupDf.printSchema()
```

//autoBroadcast Join is turned off

```
spark.sql("SET spark.sql.autoBroadcastJoinThreshold = -1")
```

```
val joinCondition = batsmenBestRunsAvgHistoryDf.col("Batsman") ===
batsmenWorldCupDf.col("batsman")
```

```
val joinType = "inner"
```

//Using broadcast join

```
val finalBestBatsmenPlayingWorldCupDf =
batsmenBestRunsAvgHistoryDf.join(broadcast(batsmenWorldCupDf),joinCondition,joinT
ype).drop (batsmenBestRunsAvgHistoryDf.col("Batsman"))
```

```
finalBestBatsmenPlayingWorldCupDf.orderBy(desc("AverageRunsScored")).show()
```

```
spark.stop()
```

```
}
```

Output:

```
+-----+-----+
|AverageRunsScored|    batsman|
```

111.5	Rohit_Sharma
100.0	Virat_Kohli
77.0	Steven_Smith

