

WEEK 16 QUIZ SOLUTION

1. Choose the odd one out. Normal spark streaming takes care of:
 - a. automatic looping between the batches
 - b. storing intermediate results
 - c. combining results to previous batch result
 - d. ***Future upgrades and enhancements**

Explanation: D is odd man out because future enhancements are expected, its not provided right now.

2. Normal spark streaming (Dstream) is built on top of RDDs
 - a. ***TRUE**
 - b. FALSE
3. Spark structured streaming API:
 - a. offers unified model for batch and stream processing.
 - b. runs over the spark sql engine. Takes advantage of the sql based optimizations.
 - c. native support for event time semantics.
 - d. ***All of the above**

Explanation: All of the above are true in regards to Spark structured streaming API

4. Append output mode is not supported with streaming aggregation
 - a. ***TRUE**
 - b. FALSE
5. Spark structure Streaming supports spark sql engine optimization.
 - a. ***TRUE**
 - b. FALSE
6. How many shuffle partitions will be created by default?
 - a. 2
 - b. 20
 - c. ***200**
 - d. 2000

WEEK 16 QUIZ SOLUTION

Explanation: In streaming, default number of partitions is 200

7. Which DataStream writer output mode is not allowed while doing a streaming aggregation?

- a. update
- b. Insert
- c. ***append**
- d. complete

Explanation: Append is not support for streaming aggregation.

8. Choose the correct configuration to gracefully shutdown Spark streaming application.

- a. spark.streaming.config("stopGracefullyOnShutdown","false")
- b. spark.streaming.config("stop.GracefullyOnShutdown","true")
- c. ***.config("spark.streaming.stopGracefullyOnShutdown","true")**
- d. .config("spark.streaming.stop.GracefullyOnShutdown","true")

Explanation:

`.config("spark.streaming.stopGracefullyOnShutdown","true")`

9. Data Stream reader allows us to define the trigger.

- a. TRUE
- b. ***FALSE**
- c. Explanation: Data Stream write supports defining the trigger.

10. A series of fixed size, overlapping window is called:

- a. slicing time window
- b. Tumbling time window
- c. Overlapping time window
- d. ***sliding time window**

Explanation: A series of fixed size, overlapping window is called sliding time window

WEEK 16 QUIZ SOLUTION

11. A _____ is a threshold to specify how long the system waits for late events
- a. ***watermark**
 - b. sliding window
 - c. batch interval
 - d. event time

Explanation: A **watermark** is a threshold to specify how long the system waits for late events

12. The default output mode of a spark streaming event?
- a. ***append**
 - b. update
 - c. complete
 - d. insert

Explanation: append is default output mode of a spark streaming event

13. Events within the watermark may or may not be taken.
- a. TRUE
 - b. ***FALSE**

Explanation: Events within the watermark are considered

14. Stateful operations can lead to out of memory error.
- a. ***TRUE**
 - b. FALSE

15. Which built-in output sink stores the output to a specified directory?
- a. Console Sink
 - b. Memory Sink
 - c. ***File Sink**
 - d. Foreach Sink

Explanation: File is a built in stores the output to a specified directory.